

# Dynamic Resource Management using Operating System-Level Virtualization

Computer Science 4490z  
Undergraduate Research Project Thesis

Alexander Pokluda

Department of Computer Science  
University of Western Ontario

April 10, 2010

# Outline

- 1 Introduction
- 2 OpenVZ
  - Resource Management
  - Managing Resource Stress Situations
- 3 System Overview
  - Architecture Overview
  - Design Overview
  - Source Code Availability and Documentation
- 4 Validation
  - Test Environment
  - Experimental Results
- 5 Conclusion
  - Future Work
  - Summary

# Outline

## 1 Introduction

## 2 OpenVZ

- Resource Management
- Managing Resource Stress Situations

## 3 System Overview

- Architecture Overview
- Design Overview
- Source Code Availability and Documentation

## 4 Validation

- Test Environment
- Experimental Results

## 5 Conclusion

- Future Work
- Summary

# Outline

- 1 Introduction
- 2 OpenVZ
  - Resource Management
  - Managing Resource Stress Situations
- 3 System Overview
  - Architecture Overview
  - Design Overview
  - Source Code Availability and Documentation
- 4 Validation
  - Test Environment
  - Experimental Results
- 5 Conclusion
  - Future Work
  - Summary

# Outline

- 1 Introduction
- 2 OpenVZ
  - Resource Management
  - Managing Resource Stress Situations
- 3 System Overview
  - Architecture Overview
  - Design Overview
  - Source Code Availability and Documentation
- 4 Validation
  - Test Environment
  - Experimental Results
- 5 Conclusion
  - Future Work
  - Summary

# Outline

- 1 Introduction
- 2 OpenVZ
  - Resource Management
  - Managing Resource Stress Situations
- 3 System Overview
  - Architecture Overview
  - Design Overview
  - Source Code Availability and Documentation
- 4 Validation
  - Test Environment
  - Experimental Results
- 5 Conclusion
  - Future Work
  - Summary

# Outline

- 1 Introduction
- 2 OpenVZ
  - Resource Management
  - Managing Resource Stress Situations
- 3 System Overview
  - Architecture Overview
  - Design Overview
  - Source Code Availability and Documentation
- 4 Validation
  - Test Environment
  - Experimental Results
- 5 Conclusion
  - Future Work
  - Summary

# Types of Virtualization

## Definition

Virtualization is the abstraction of computer resources

- **Platform virtualization** enables the execution of one or more **virtual machines** on a single computer
- Modern types of platform virtualization include:
  - *full virtualization*
  - *hardware-assisted virtualization*
  - *paravirtualization*
  - *operating system-level virtualization*
- In operating system-level virtualization, all virtual machines (also known as **virtual environments**) share one operating system kernel
- Operating system-level virtualization has very little overhead so applications can achieve near-native performance

# Motivation for Using Virtualization

- By using virtual machines to run many independent software systems on a single physical server, greater resource utilization levels can be achieved
- Greater resource utilization levels mean that less physical resources are required and overall costs are reduced

# Problem Statement

## New Problem

How do we effectively manage the resources of a cluster of hardware nodes as a single unit?

- This thesis expands upon a system called **Golondrina**
- Golondrina works by identifying localized **resource stress** situations then attempting to dissipate them by reallocating system resources and, if necessary, by **migrating or replicating** virtual environments

## Contributions:

- memory management studied
- heuristic developed
- new architecture
- CPU code as plug-in
- memory plug-in developed
- documentation
- testing

# Outline

## 1 Introduction

## 2 OpenVZ

- Resource Management
- Managing Resource Stress Situations

## 3 System Overview

- Architecture Overview
- Design Overview
- Source Code Availability and Documentation

## 4 Validation

- Test Environment
- Experimental Results

## 5 Conclusion

- Future Work
- Summary

# OpenVZ

## What is OpenVZ?

**OpenVZ** is a mature open-source community project that implements operating system-level virtualization using Linux

OpenVZ provides four primary controls for per-container resource accounting and limiting:

- **user bean counters**
- **disk quota management**
- **CPU fair scheduler**
- **configurable input/output priorities**

System administrators can also use standard Linux resource management and accounting mechanisms such as `tc` and `iptables`

# Memory Allocation

- OS-level virtualization makes it trivial to add or remove arbitrary amounts of memory to or from a container in real time. In OpenVZ, this is done with the user beancounters.
- There are currently 24 beancounters and each has 5 attributes: held, maxheld, barrier, limit, failcnt

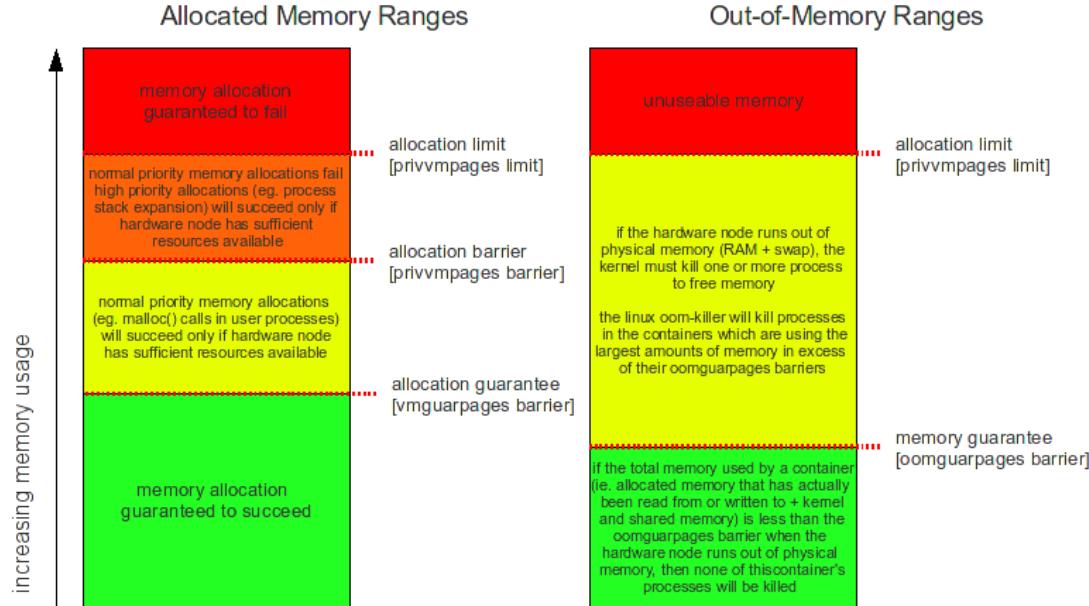
Three main user beancounters relate directly to memory management:

`vmguarpages` guaranteed virtual memory pages

`privvmpages` private virtual memory pages

`oomguarpages` out-of-memory guaranteed memory pages

- In Linux and OpenVZ memory is overcommitted by default



**Figure: Typical relative settings of the privvmpages, vmguarpages and oomguarpages user bean counters for a container**

# Golondrina and OpenVZ Resource Management

## First Version

- monitored a single resource: the processor
- monitored CPU time using the virtual file `/proc/vz/vestat`
- dissipated resource stress situations using **migrations** or **replications**
- has been packaged as a plug-in in the current version

## Current Version

- memory management plug-in developed
- monitors memory usage within each container and for hardware node using `user_beancounters` virtual file and `vzmemcheck`
- dissipates resource stress using **user beancounters**, **migrations** and **replications**

# Managing Resource Stress Situations

- Goal: ensure resource availability for each virtual machine while achieving high levels of resource utilization on each hardware node
- Ideally containers should be distributed such that a predetermined level of resource utilization, say  $x$ , is achieved on each hardware node
- This is equivalent to the NP-complete subset-sum problem

## Proof.

Let  $S_0$  be the set of containers to be distributed across several hardware nodes. If we wish to achieve a target resource utilization level of  $x$  on each hardware node, then we need to find a subset  $s_0 \subseteq S_0$ , such that the sum of the resource usage of each container in  $s_0$  is exactly  $x$ . Once this subset has been found, the containers in  $s_0$  are assigned to a hardware node and the process repeats for  $S_1 = S_0 - s_0$ .



- Finding good resource allocation strategies is challenging and an active area of research

# Resource Stress Indicators for Memory

## ① An increase in oomguarpages failcnt value

- **raw score:** the increase in oomguarpages failcnt, i.e. the number of processes that have been killed
- **normalized score** =  $\begin{cases} 0.0 & \text{if rawscore is 0} \\ 1.0 & \text{otherwise} \end{cases}$

## ② An increase in privvmpages failcnt value

- **raw score:** the increase in privvmpages failcnt, i.e. the number of failed memory allocation attempts
- **normalized score** =  $\begin{cases} 0.0 & \text{if rawscore is 0} \\ 1.0 & \text{otherwise} \end{cases}$

## ③ Current memory usage (oomguarpages held + kmemsize held + \*buf held) versus oomguarpages barrier

- **raw score:**  $(\text{oomguarpages held} + \text{kmemsize held} + \text{*buf held}) / \text{oomguarpages barrier}$ , i.e. the fraction of guaranteed memory used
- **normalized score** =  $\min\{1.0, \text{rawscore}\}$

## ④ privvmpages held versus privvmpages barrier

- **raw score:**  $\text{privvmpages held} / \text{privvmpages barrier}$ , i.e. the fraction of memory the memory allocation guarantee used
- **normalized score** =  $\min\{1.0, \text{rawscore}\}$

# Resource Stress Resolution for Memory

A basic heuristic policy for resolving a memory stress situation in a container:

- ① Increase the memory limit for the stressed container
- ② Migrate stressed container to another hardware node
- ③ Migrate another container
- ④ Alternatively, start a replica

A basic heuristic policy for resolving a memory stress situation on a hardware node:

- ① Migrate the container that is using the largest amount of resources
- ② Migrate the container that is using next largest amount of resources
- ③ Repeat

# Outline

## 1 Introduction

## 2 OpenVZ

- Resource Management
- Managing Resource Stress Situations

## 3 System Overview

- Architecture Overview
- Design Overview
- Source Code Availability and Documentation

## 4 Validation

- Test Environment
- Experimental Results

## 5 Conclusion

- Future Work
- Summary

# Architecture Overview

- The system has been architected as a research system that will be a test bed for various resource management policies
- The primary architectural pattern for Golondrina is “client-server”

## Client Component

- collects resource usage statistics and sends them to the server

## Gate Component

- manages configuration for an external load balancer

## Server Component

- Analyses resource usage statistics to identify resource stress situations
- Instructs client components to adjust resource limits and perform migrations and replications
- Instructs gate component to update load balancing configuration

# Architecture Overview

- The system has been architected as a research system that will be a test bed for various resource management policies
- The primary architectural pattern for Golondrina is “client-server”

## Client Component

- collects resource usage statistics and sends them to the server

## Gate Component

- manages configuration for an external load balancer

## Server Component

- Analyses resource usage statistics to identify resource stress situations
- Instructs client components to adjust resource limits and perform migrations and replications
- Instructs gate component to update load balancing configuration

# Architecture Overview

- The system has been architected as a research system that will be a test bed for various resource management policies
- The primary architectural pattern for Golondrina is “client-server”

## Client Component

- collects resource usage statistics and sends them to the server

## Gate Component

- manages configuration for an external load balancer

## Server Component

- Analyses resource usage statistics to identify resource stress situations
- Instructs client components to adjust resource limits and perform migrations and replications
- Instructs gate component to update load balancing configuration

# Architecture Overview

- The system has been architected as a research system that will be a test bed for various resource management policies
- The primary architectural pattern for Golondrina is “client-server”

## Client Component

- collects resource usage statistics and sends them to the server

## Gate Component

- manages configuration for an external load balancer

## Server Component

- Analyses resource usage statistics to identify resource stress situations
- Instructs client components to adjust resource limits and perform migrations and replications
- Instructs gate component to update load balancing configuration

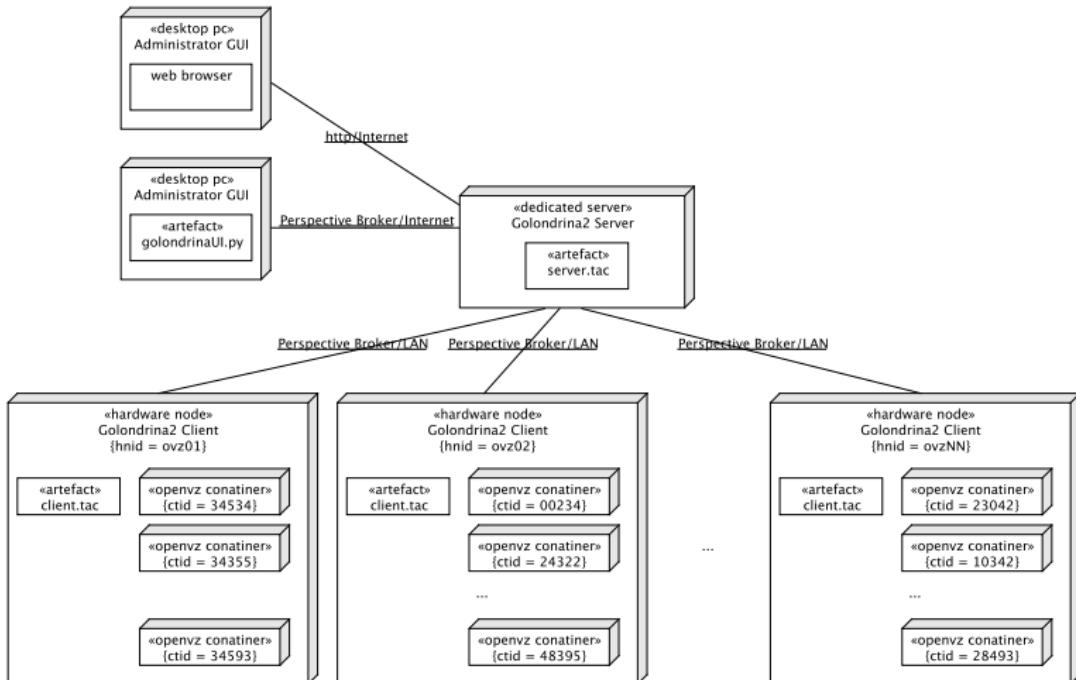


Figure: Golondrina deployment diagram

# Design Overview

- The software for collecting and analysing resource usage statistics has been encapsulated in **plug-ins**
  - sensor plug-ins
  - overload identifier plug-ins
  - overload resolver plug-ins
- Each overload resolver plug-in is given a priority value
- Once a resource stress has been identified, the overload resolver plug-ins for the stressed resource(s) are each given a chance in turn to try and dissipate the resource stress
- Certain aspects of a policy's state (e.g. a threshold value) can be updated at runtime
- A configuration manager subsystem maintains a database of configuration information for the whole system and can be used to alter the system's run-time behaviour

# Source Code Availability and Documentation

- The source code developed as part of this thesis is free software, anyone is free to redistribute it and/or modify it under the terms of the GNU General Public Licence
- Source code and documentation is available at:  
<http://alexanderpokluda.ca/trac/cs4490>



The screenshot shows a Mozilla Firefox window with the title bar "Dynamic Resource Management using OS Virtualization - Mozilla Firefox". The address bar contains the URL <http://alexanderpokluda.ca/trac/cs4490>. The main content area displays the Trac project page for "Dynamic Resource Management with OS Virtualization". The page features a header with the project name in blue and grey, and navigation links for Login, Preferences, Help/Guide, About Trac, and My Notifications. Below the header are links for Wiki, Timeline, Roadmap, Browse Source, View Tickets, Search, and API Docs. A secondary navigation bar at the bottom includes Start Page, Index, History, Last Change, and Watch Page. The main content area contains a brief introduction to the Golondrina system and a Table of Contents sidebar with links to Download, Source Checkout, System Requirements, Setting up Eclipse, Terminology, and OpenVZ Resource Management.

This project expands upon a system called Golondrina. Golondrina is a system that performs autonomic workload management among a cluster of hardware nodes running OpenVZ. Golondrina works by identifying localized resource stress situations then

Table of Contents

[Download](#)  
[Source Checkout](#)  
[System Requirements](#)  
[Setting up Eclipse](#)  
[Terminology](#)  
[OpenVZ Resource Management](#)

# Outline

## 1 Introduction

## 2 OpenVZ

- Resource Management
- Managing Resource Stress Situations

## 3 System Overview

- Architecture Overview
- Design Overview
- Source Code Availability and Documentation

## 4 Validation

- Test Environment
- Experimental Results

## 5 Conclusion

- Future Work
- Summary

# Test Environment

Two configurations were investigated for the virtual machine to be placed under load:

- TPC-W benchmark
- LAMP software stack
- Joomla! was selected as the specific application to be tested running on the LAMP stack
- Apache JMeter was used for load generation and performance monitoring
- Containers created using CentOS 5.3 template
- Also installed: Apache HTTP Server ver. 2.2.3, MySQL ver. 5.0.45, PHP ver. 5.1.6, and Joomla! ver. 1.5.14
- Three identical hardware nodes, each with 3.40 GHz dual-core processor, 2 GiB RAM, 2 GiB swap
- Apache HTTP Server **prefork** module used

# Experimental Results

- Test 0 – No Memory Stress
- Test 1 – Unresolved Memory Stress
- Test 2 – Memory Stress Resolved Locally
- Test 3 – Memory Stress Resolved with Migration

Test	Avg	Min	Max	Std Dev	Err %	Throughput	Fail Count
0	448	166	2275	243.72	0.00	8.0	0
1	1082	124	55576	5817.40	2.22	3.6	1002
2	866	115	49214	4748.84	1.63	4.9	809
3	370	137	4354	419.42	9.24	5.5	1143

Table: Results for Four Tests, each with four runs

# Outline

## 1 Introduction

## 2 OpenVZ

- Resource Management
- Managing Resource Stress Situations

## 3 System Overview

- Architecture Overview
- Design Overview
- Source Code Availability and Documentation

## 4 Validation

- Test Environment
- Experimental Results

## 5 Conclusion

- Future Work
- Summary

# Future Work

- Experiments involving memory stresses and replication
- Study the interaction between policies when a container is experiencing a stress for more than one resource

Possible improvements in future versions:

- Add a mechanism to reclaim resources
- Use remote storage for container private areas
- Distribute decision making
- Add memory resource usage prediction

# Summary

- Golondrina is a system that performs dynamic resource management among a cluster of hardware nodes
- Different models of virtualization and the advantages of operating system-level virtualization were discussed
- The heuristic currently used to detect and dissipate memory stress situations was presented
- The architecture and design of Golondrina was discussed
- The functionality of the system was validated using a series of experimental tests
- A brief summary of experiments yet to be performed and possible future enhancements was presented